

COMP9334


# Capacity Planning for Computer Systems and Networks

---


## Week 3B: Markov Chain

# Last lecture: Queues with Poisson arrivals

- Single-server




- Multi-server



# This week: Markov Chain


- You can use Markov Chain to analyse
  - Closed queueing network (see example below)
  - Reliability problem



- There are  $n$  jobs in the closed system
- What is the response time of one job?
- What is the response time if we replace the CPU with one that is twice as fast?

# Markov chain

- The state-transition model that we have used is called a continuous-time Markov chain
  - There is also discrete-time Markov chain
- The transition from a state of the Markov chain to another state is characterised by an exponential distribution
  - E.g. The transition from State  $p$  to State  $q$  is exponential with rate  $r_{pq}$ , then consider a small time interval  $\delta$
  - Prob [ Transition from State  $p$  to State  $q$  in time  $\delta$  | State  $p$ ] =  $r_{pq} \delta$




# Method for solving Markov chain

---

- A Markov chain can be solved by
  - Identifying the states
  - Find the transition rate between the states
  - Solve the steady state probabilities
- You can then use the steady state probabilities as a stepping stone to find the quantity of interest (e.g. response time etc.)
- We will study two Markov chain problems in this lecture:
  - Problem 1: A Database server
  - Problem 2: Data centre reliability problem

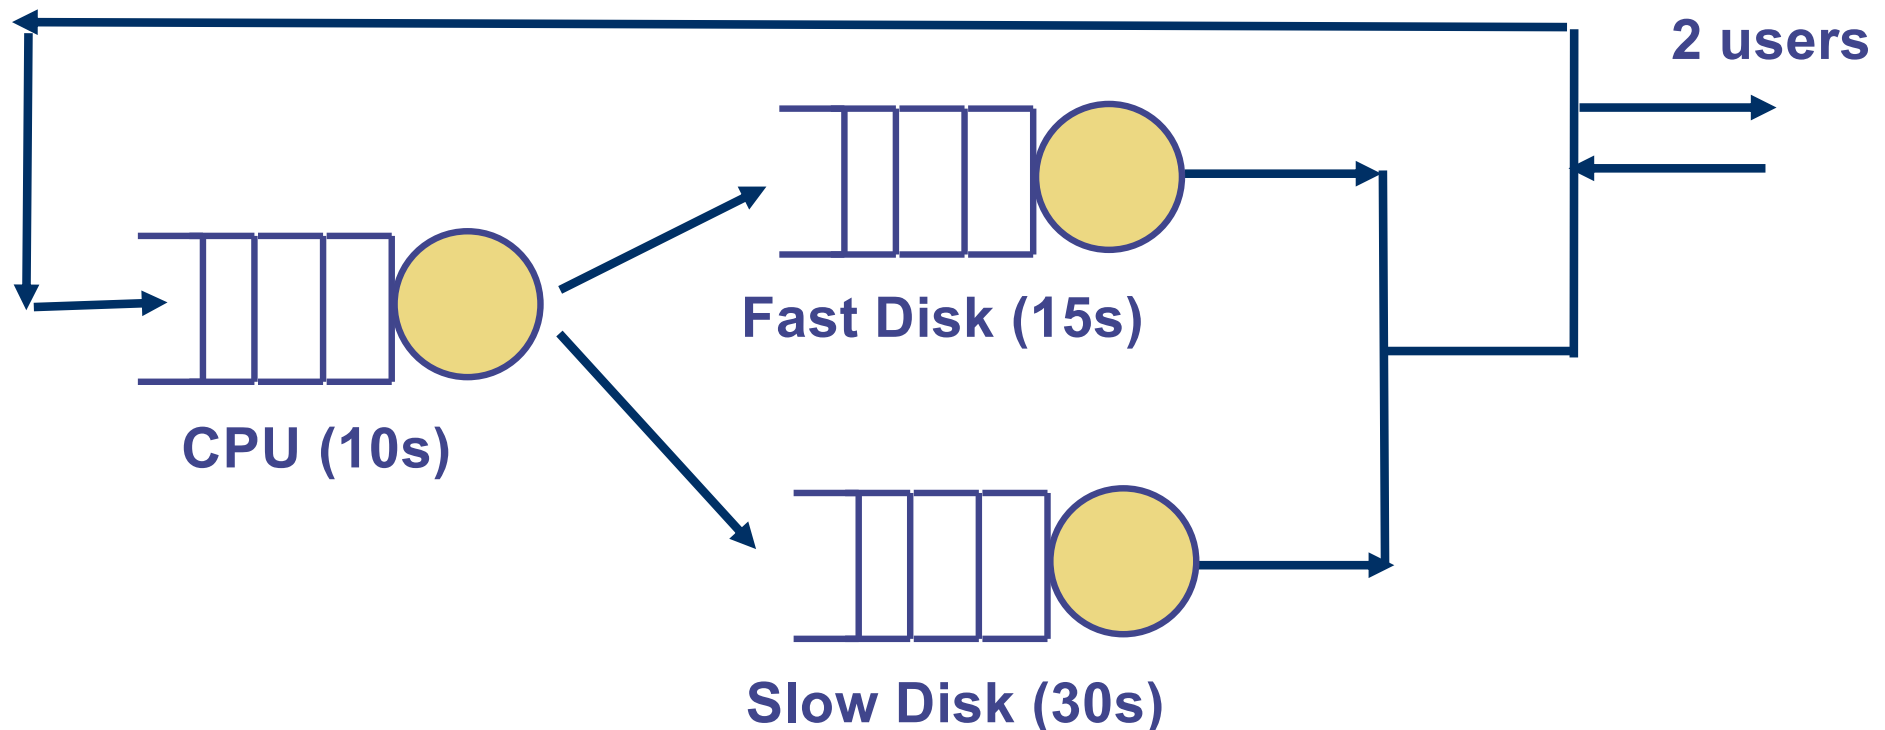
# Problem 1: A DB server

- A database server with a CPU, a fast disk and a slow disk
- At peak demand, there are always two users in the system
- Users alternate between the CPU and the disks
- The users will equally likely find the file on either disk



## Problem 1: A DB server (cont'd)

- Fast disk is twice as fast as the slow disk
- Typical transactions take on average 10s CPU time
- Fast disk takes on average 15s to serve all files for a user
- Slow disk takes on average 30s to serve all files for a user
- The time that each transaction requires from the CPU and the disks is exponentially distributed



# Typical capacity planning questions

---

- What response time can a typical user expect?
- What is the utilisation of each of the system resources?
- How will performance parameters change if number of users are doubled?
- If fast disk fails and all files are moved to slow disk, what will be the new response time?




# Choice of states #1

---

- Use a 2-tuple (A,B) where
  - A is the location of the first user
  - B is the location of the second user
  - A, B are drawn from {CPU,FD,SD}
    - FD = fast disk, SD = slow disk
  - Example states are:
    - (CPU,CPU): both users at CPU
    - (CPU, FD): 1st user at CPU, 2nd user at fast disk
  - Total 9 states
- Question: If there are  $n$  users,
  - What are the states?
  - How many states are there?

## Choice of states #2

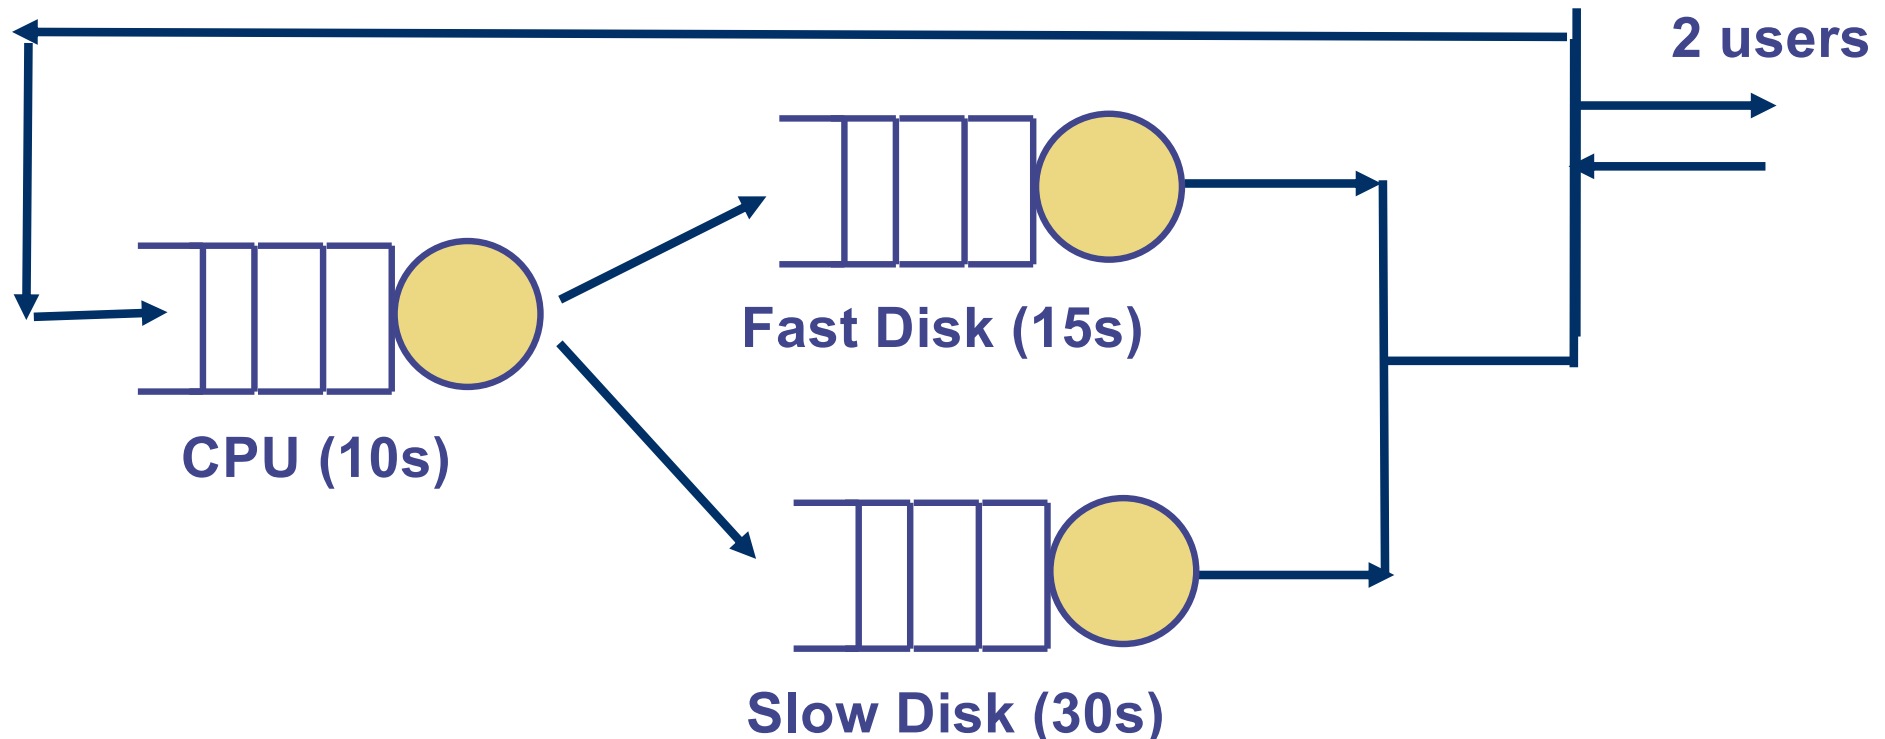
- We use a 3-tuple (X,Y,Z)
  - X is # users at CPU
  - Y is # users at fast disk
  - Z is # users at slow disk
- Examples
  - (2,0,0): both users at CPU
  - (1,0,1): one user at CPU and one user at slow disk
- There are six possible states. Can you list them?
  - 
- If there are n users, how many states are there?

$$\frac{(n+1)(n+2)}{2}$$

Choice #2 requires less #states but loses certain information.


# Identifying state transitions (1)

- A state is: (#users at CPU, #users at fast disk, #users at slow disk)
- What is the rate of moving from State (2,0,0) to State (1,1,0)?
  - This is caused by a job finishing at the CPU and move to fast disk
  - Jobs complete at CPU at a rate of 6 transactions/minute
  - Half of the jobs go to the fast disk
- Transition rate from (2,0,0)  $\rightarrow$  (1,1,0) = 3 transactions/minute
- Similarly, transition rate from (2,0,0)  $\rightarrow$  (1,0,1) = 3 transactions/minute



## State transition diagram (2)

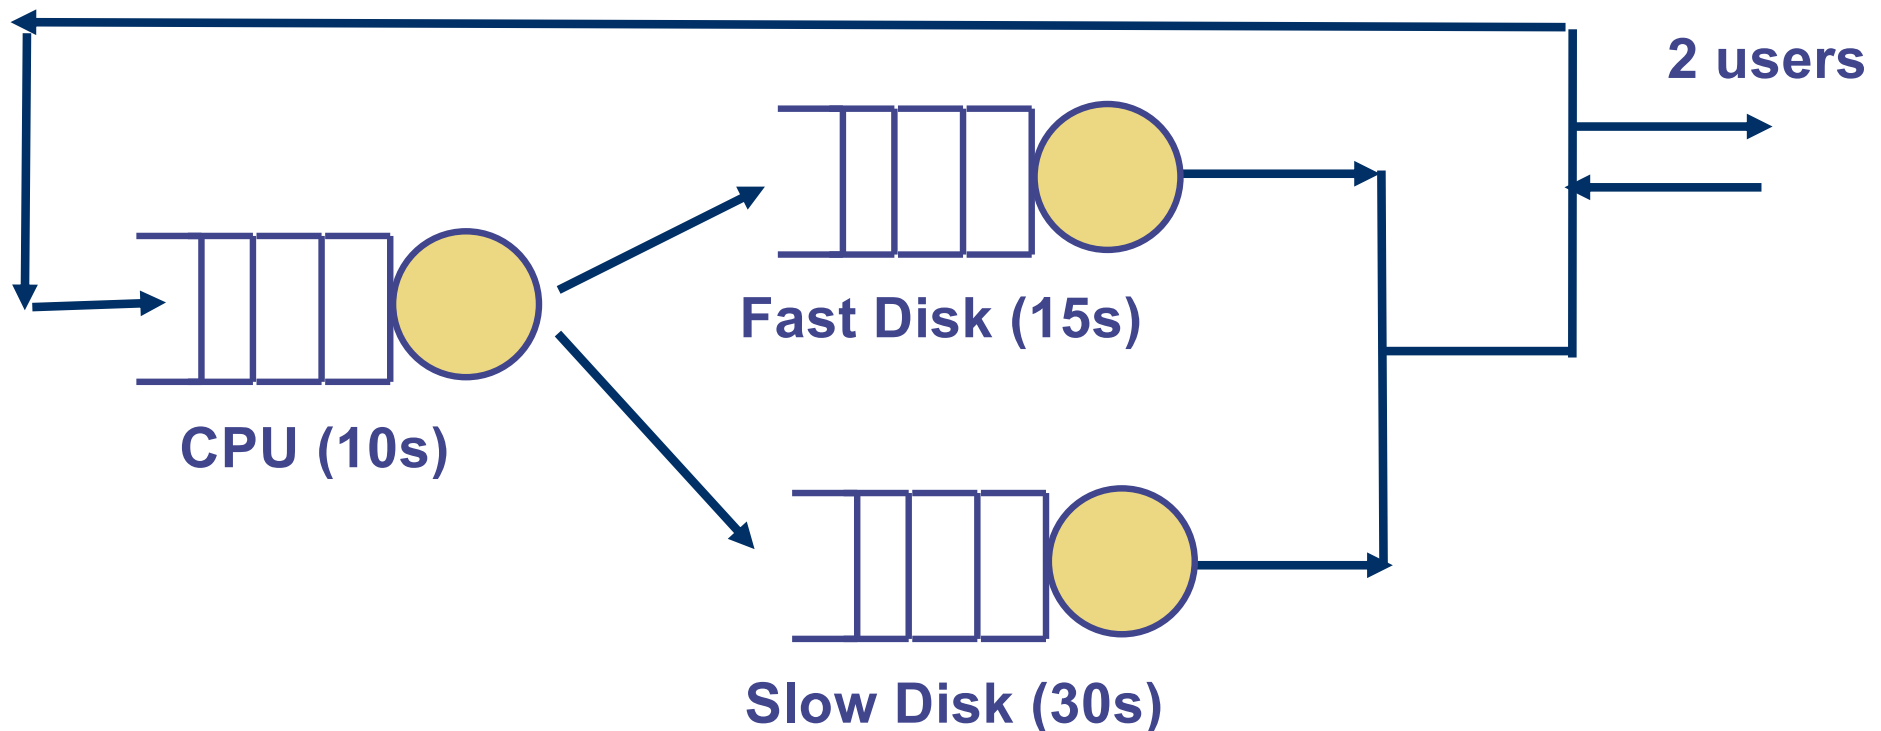
- Transition rate from  $(2,0,0) \rightarrow (1,1,0) = 3$  transactions/minute
- Transition rate from  $(2,0,0) \rightarrow (1,0,1) = 3$  transactions/minute




- Question: What is the transition rate from  $(2,0,0) \rightarrow (0,1,1)$ ?

## Identifying state transitions (2)

- From (1,1,0) there are 3 possible transitions
  - Fast disk user goes back to CPU (2,0,0)
  - CPU user goes to the fast disk (0,2,0), or
  - CPU user goes to the slow disk (0,1,1)
- Question: What are the transition rates in number of transactions per minute?



# Completing the state transition diagram



## Exercise

---

- The state transition diagram is still not complete. Choose any two state transitions and determine their rates.

# Complete state transition diagram

---






# Balance Equations

Define

$P_{(2,0,0)}$  = Probability in state (2,0,0)

$P_{(1,1,0)}$  = Probability in state (1,1,0) etc.

Exercise: Write down the balance equation for state (2,0,0)



## Flow balance equations

- You can write one flow balance equation for each state:

$$6 P_{(2,0,0)} - 4 P_{(1,1,0)} - 2 P_{(1,0,1)} + 0 P_{(0,2,0)} + 0 P_{(0,1,1)} + 0 P_{(0,0,2)} = 0$$

$$-3 P_{(2,0,0)} + 10 P_{(1,1,0)} + 0 P_{(1,0,1)} - 4 P_{(0,2,0)} - 2 P_{(0,1,1)} + 0 P_{(0,0,2)} = 0$$

$$-3 P_{(2,0,0)} + 0 P_{(1,1,0)} + 8 P_{(1,0,1)} + 0 P_{(0,2,0)} - 4 P_{(0,1,1)} - 2 P_{(0,0,2)} = 0$$

$$0 P_{(2,0,0)} - 3 P_{(1,1,0)} + 0 P_{(1,0,1)} + 4 P_{(0,2,0)} + 0 P_{(0,1,1)} + 0 P_{(0,0,2)} = 0$$

$$0 P_{(2,0,0)} - 3 P_{(1,1,0)} - 3 P_{(1,0,1)} + 0 P_{(0,2,0)} + 6 P_{(0,1,1)} + 0 P_{(0,0,2)} = 0$$

$$0 P_{(2,0,0)} + 0 P_{(1,1,0)} - 3 P_{(1,0,1)} + 0 P_{(0,2,0)} + 0 P_{(0,1,1)} + 2 P_{(0,0,2)} = 0$$

- However, there are only 5 linearly independent equations.
- Need one more equation:

# Steady State Probability

---

- You can find the steady state probabilities from 6 equations
  - It's easier to solve the equations by a software packages, e.g
    - Python, Matlab, Octave, etc.
- The solutions are:
  - $P_{(2,0,0)} = 0.1391$
  - $P_{(1,1,0)} = 0.1043$
  - $P_{(1,0,1)} = 0.2087$
  - $P_{(0,2,0)} = 0.0783$
  - $P_{(0,1,1)} = 0.1565$
  - $P_{(0,0,2)} = 0.3131$
- I used Python (the numpy library) to solve these equations
  - The file is “data\_server.py” (can be downloaded from the course web site)
- How can we use these results for capacity planning?

# Model interpretation

- Response time of each transaction
  - Use Little's Law  $R = N/X$  with  $N = 2$ 
    - For this system:
      - System throughput = CPU Throughput
    - Throughput = Utilisation x Service **rate**
      - Recall Utilisation = Throughput x Service **time** (From Lecture 1B)
    - CPU utilisation (using states where there is a job at CPU):  
 $P_{(2,0,0)} + P_{(1,1,0)} + P_{(1,0,1)} = 0.452$
    - Throughput =  $0.452 \times 6 = 2.7130$  transactions / minute
    - Response time (with 2 users) =  $2 / 2.7126 = 0.7372$  minutes per transaction

# Sample capacity planning problem

---

- What is the response time if the system has up to 4 users instead of 2 users only?
  - You can't use the previous Markov chain
  - You need to develop a new Markov chain
    - The states are again (#users at CPU, #users at fast disk, #users at slow disk)
    - States are (4,0,0), (3,1,0), (1,2,1) etc.
    - There are 15 states
    - Determine the transition rates
    - Write down the balance equations and solve them.
    - Use the steady state probabilities and Little's Law to determine the new response time
    - You can do this as an exercise
    - Throughput = 3.4768 (up 28%), response time = 60.03 seconds (up 56%)


# Computation aspect of Markov chain

---

- This example shows that when there are a large number of users, the burden to build a Markov chain model is large
  - 15 states
  - Many transitions
  - Need to solve 15 equations in 15 unknowns
- Is there a faster way to do this?
  - Yes, we will look at Mean Value Analysis in a few weeks and it can obtain the response time much more quickly

# Machine working-repair cycle

- A data centre consists of a number of machines
- Machines can fail and have to be repaired
- Terminology:
  - Time-to-next failure: From the time a machine has been fixed to the time it next fails
  - Time-to-repair = time waiting in the repair queue + service time to repair the machine
    - Time-to-repair is a response time



# Data centre reliability problem

---

- Example: A data centre has 10 machines
  - Each machine may go down
    - Time-to-next-failure is exponentially distributed with mean 90 days
  - Service time to repair is exponentially distributed with mean 6 hours
- Capacity planning question:
  - Can I make sure that at least 8 machines are available 99.9999% of the time?
  - What is the probability that at least 6 machines are available?
  - How many repair staff are required to guarantee that at least  $k$  machines are available with a given probability?
  - What is the mean-time-to-repair (MTTR) a machine?
    - Note: Mean-time-to-repair includes waiting time at the repair queue.




# Data centre reliability - general problem

---

- Data centre has
  - $M$  machines
  - $N$  staff maintain and repair machine
  - Assumption:  $M > N$
- Automatic diagnostic system
  - Check “heartbeat” by “ping” (Failure detection)
  - Staff are informed if failure is detected
- Repair work
  - If a machine fails, any one of the idle repair staff (if there is one) will attend to it.
  - If all repair staff are busy, a failed machine will need to wait until a repair staff has finished its work
- This is a queueing problem solvable by Markov chain!!!
- Let us denote
  - $\lambda = 1 / \text{Mean-time-to-failure}$
  - $\mu = 1 / \text{Mean service time to repair a machine}$

# Queueing model for data centre example

An arrival is due to a machine failure.




A departure occurs when a machine has been repaired.

We build a Markov chain for this box.

## Markov model for the repair queue


- State  $k$  represents  $k$  machines have failed
- Part of the state transition diagram is showed below



The rate of failure for one machine is  $\lambda$ . In State 0, there are  $M$  working machine, the failure rate is  $M\lambda$ .

The same argument holds for other state transition probability.

# Markov Model for the repair queue



Note: There are  $(M+1)$  states.

Why is it  $N\mu$ ?

Why not  $(N+1)\mu$ ?

## Solving the model



- We can solve for  $P(0)$ ,  $P(1)$ , ...,  $P(M)$

$$P(k) = \begin{cases} P(0) \left(\frac{\lambda}{\mu}\right)^k C_k^m & k = 1, \dots, N \\ P(0) \left(\frac{\lambda}{\mu}\right)^k C_k^m \frac{N^{N-k} k!}{N!} & k = N + 1, \dots, M \end{cases}$$

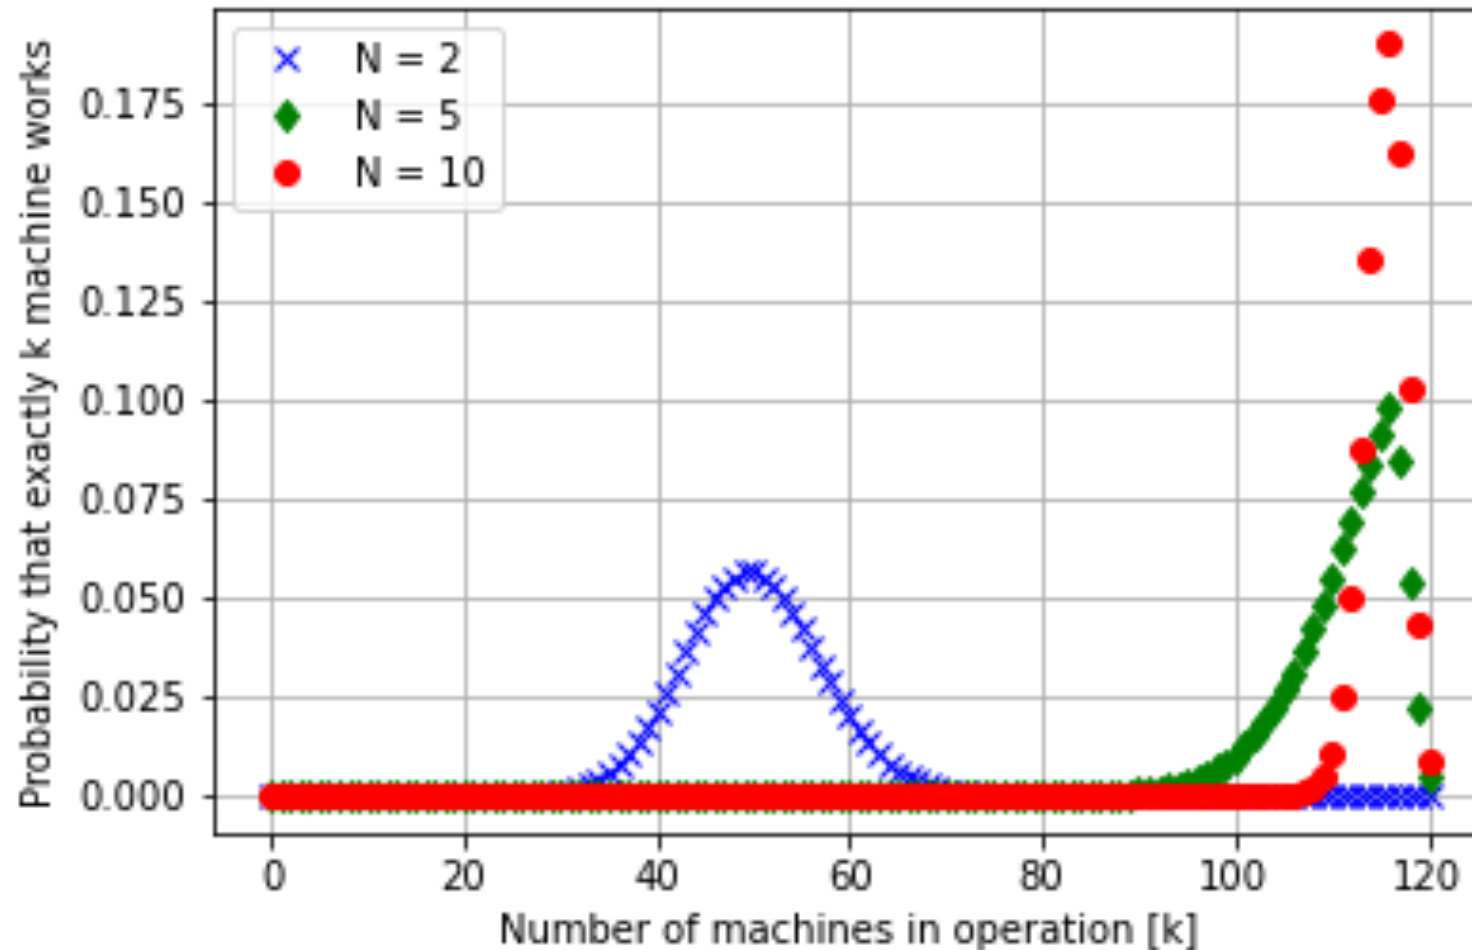
Where

$$P(0) = \left[ \sum_{k=0}^N \left(\frac{\lambda}{\mu}\right)^k C_k^m + \sum_{k=N+1}^M \left(\frac{\lambda}{\mu}\right)^k C_k^m \frac{N^{N-k} k!}{N!} \right]^{-1}$$


# Using the model

- Probability that exactly  $k$  machines are available = 
- Probability that at least  $k$  machines are available  
= 
- But expression for  $P(k)$ 's are complicated, need numerical software
- Example:
  - $M = 120$
  - Mean-time-to-failure = 500 minutes
  - Mean service time to repair = 20 minutes
  - $N = 2, 5$  or  $10$
  - The results are showed in the graphs in the next 2 pages
    - I used the file “data\_centre.py” to do the computation, the file is available on the course web site.

# Probability that exactly $k$ machines operate



# Probability that at least $k$ machines operate






Think time ~ Mean-time-to-failure (MTTF) =  $1 / \lambda$

Throughput  
~ Mean machine failure  
rate  
(see next page)

Mean time to repair  
(MTTR)  
= Queueing time for  
repair + actual repair  
time

Can compute MTTR  
using Little's Law.



## Mean machine failure rate

State	Probability	Failure rate
0	$P(0)$	$M\lambda$
1	$P(1)$	$(M-1)\lambda$
2	$P(2)$	$(M-2)\lambda$
...	...	
k	$P(k)$	$(M-k)\lambda$
...	...	
M	$P(M)$	0

$$\bar{X}_f = \sum_{k=0}^{M-1} (M - k)\lambda P(k)$$

# Continuous-time Markov chain

---

- Useful for analysing queues when the inter-arrival or service time distribution is exponential
- The procedure is fairly standard for obtaining the steady state probability distribution
  - Identify the state
  - Find the state transition rates
  - Set up the balance equations
  - Solve the steady state probability
- We can use the steady state probability to obtain other performance metrics: throughput, response time etc.
  - May need Little's Law etc.
- Continuous-time Markov chain is only applicable when the underlying probability distribution is exponential but the operations laws (e.g. Little's Law) are applicable no matter what the underlying probability distributions are.

# Markov chain

---

- Markov chain is big field in itself. We have touched on only continuous-time Markov chain
  - There are also discrete time Markov chains
  - Markov chain has discrete state, a related concept is Markov process whose states are continuous
- Markov chain / processes have many applications
  - Page rank algorithm from Google can be explained in terms of discrete-time Markov chain
  - Graphical Models (from machine learning)
  - Transport engineering
  - Mathematical finance
- Personally, I use Markov chains to understand how living cells process information

# References

---

- Recommended reading
  - The database server example is taken from Menasce et al., “Performance by design”, Chapter 10
  - The data centre example is taken from Mensace et al, “Performance by design”, Chapter 7, Sections 1-4
- For a more in-depth, and mathematical discussion of continuous-time Markov chain, see
  - Alberto Leon-Gracia, “Probabilities and random processes for Electrical Engineering”, Chapter 8.
  - Leonard Kleinrock, “Queueing Systems”, Volume 1